







DPO | 90% previous reference | Measured policy probabilities
alpha=1, nu=0.9, lambda=0.8, beta=0.8, seed=0
 $r_t = 0s_0 + 0.1s_t + 0.9s_{t-1}$ | Sequence-sum, normalized within four responses



— Response 1 - - Response 2 - . - Response 3 . . . Response 4