













IPO | 90% previous reference | Measured policy probabilities
alpha=1, nu=0.9, lambda=0.8, beta=0.2, seed=0
 $r_t = 0s_0 + 0.1s_t + 0.9s_{t-1}$ | Sequence-sum, normalized within four responses



— Response 1 - - - Response 2 - . - Response 3 . . . Response 4


DPO | 90% previous reference | Measured policy probabilities
alpha=1, nu=0.9, lambda=0.8, beta=0.8, seed=0
 $r_t = 0s_0 + 0.1s_t + 0.9s_{t-1}$ | Sequence-sum, normalized within four responses




— Response 1 - - Response 2 - . - Response 3 . . . Response 4

90% previous reference | Relative-sequence entropy
 $H(\text{softmax}(s_t - s_0))$; not raw panel entropy


Prompt 54




Prompt 251




Prompt 612




Prompt 737



Prompt 867



Prompt 945



— IPO — DPO