







DPO | Reference | Policy probabilities over iterations
alpha=0.9, lambda=0.8, beta=0.8, seed=0 | Lagged reference: nu=0.45
Sequence-sum probabilities normalized over the same four candidate responses



— Response 1 - - - Response 2 - . - Response 3 . . . Response 4