













DPO | Ordinary | $\alpha=0.9$, $\lambda=0.8$, $\beta=0.8$
Measured panel policy π : four fixed responses, all six prompts









— Response 1 - - - Response 2 - . - Response 3 . . . Response 4

DPO | Reference: nu=0.45 | alpha=0.9, lambda=0.8, beta=0.8
Measured panel policy pi: four fixed responses, all six prompts









— Response 1 - - - Response 2 - . - Response 3 . . . Response 4

DPO | Reference: nu=0.9 | alpha=0.9, lambda=0.8, beta=0.8
Measured panel policy pi: four fixed responses, all six prompts









— Response 1 - - - Response 2 - . - Response 3 . . . Response 4

DPO | Feedback: kappa=0.5 | alpha=0.9, lambda=0.8, beta=0.8
Measured panel policy pi: four fixed responses, all six prompts




— Response 1 - - - Response 2 - . - Response 3 . . . Response 4

DPO | Feedback: kappa=1 | alpha=0.9, lambda=0.8, beta=0.8
Measured panel policy pi: four fixed responses, all six prompts



— Response 1 - - - Response 2 - . - Response 3 . . . Response 4

DPO | Reference comparison | $\alpha=0.9$, $\lambda=0.8$, $\beta=0.8$
 Same prompt in each row; identical response colors and 0-1 axes





















DPO | Feedback comparison | $\alpha=0.9$, $\lambda=0.8$, $\beta=0.8$
Same prompt in each row; identical response colors and 0-1 axes

Ordinary

Feedback: $\kappa=0.5$


Feedback: $\kappa=1$




Response 1 Response 2 Response 3 Response 4

DPO | Relative-sequence entropy | alpha=0.9, lambda=0.8, beta=0.8
 $H(\text{softmax}(s(t)-s(0)))$, natural logarithm; not raw panel entropy


Prompt 54




Prompt 251




Prompt 612




Prompt 737



Prompt 867









Prompt 945









— Ordinary — Reference: nu=0.9 — Feedback: kappa=1
— Reference: nu=0.45 — Feedback: kappa=0.5

DPO | Ordinary | $\alpha=0.8$, $\lambda=0.8$, $\beta=0.8$
Measured panel policy π : four fixed responses, all six prompts









— Response 1 - - - Response 2 - . - Response 3 . . . Response 4

DPO | Reference: nu=0.4 | alpha=0.8, lambda=0.8, beta=0.8
Measured panel policy pi: four fixed responses, all six prompts









— Response 1 - - - Response 2 - . - Response 3 . . . Response 4

DPO | Reference: nu=0.8 | alpha=0.8, lambda=0.8, beta=0.8
Measured panel policy pi: four fixed responses, all six prompts









— Response 1 - - Response 2 - . Response 3 . . . Response 4

DPO | Feedback: kappa=0.5 | alpha=0.8, lambda=0.8, beta=0.8
Measured panel policy pi: four fixed responses, all six prompts




— Response 1 - - - Response 2 - . - Response 3 . . . Response 4

DPO | Feedback: kappa=1 | alpha=0.8, lambda=0.8, beta=0.8
Measured panel policy pi: four fixed responses, all six prompts




— Response 1 - - - Response 2 - . - Response 3 . . . Response 4







DPO | Reference comparison | $\alpha=0.8$, $\lambda=0.8$, $\beta=0.8$
Same prompt in each row; identical response colors and 0-1 axes



DPO | Feedback comparison | $\alpha=0.8$, $\lambda=0.8$, $\beta=0.8$
 Same prompt in each row; identical response colors and 0-1 axes









DPO | Relative-sequence entropy | alpha=0.8, lambda=0.8, beta=0.8
 $H(\text{softmax}(s(t)-s(0)))$, natural logarithm; not raw panel entropy









— Ordinary — Reference: nu=0.8 — Feedback: kappa=1
— Reference: nu=0.4 — Feedback: kappa=0.5

DPO | Ordinary | $\alpha=0.99$, $\lambda=0.8$, $\beta=0.8$
Measured panel policy π : four fixed responses, all six prompts









— Response 1 - - - Response 2 - . - Response 3 . . . Response 4

DPO | Reference: nu=0.495 | alpha=0.99, lambda=0.8, beta=0.8
Measured panel policy pi: four fixed responses, all six prompts









— Response 1 - - - Response 2 - . - Response 3 . . . Response 4

DPO | Reference: $\nu=0.99$ | $\alpha=0.99$, $\lambda=0.8$, $\beta=0.8$
Measured panel policy π : four fixed responses, all six prompts









— Response 1 - - - Response 2 - . - Response 3 . . . Response 4

DPO | Feedback: kappa=0.5 | alpha=0.99, lambda=0.8, beta=0.8
Measured panel policy pi: four fixed responses, all six prompts




— Response 1 - - - Response 2 - . - Response 3 . . . Response 4

DPO | Feedback: kappa=1 | alpha=0.99, lambda=0.8, beta=0.8
Measured panel policy pi: four fixed responses, all six prompts




— Response 1 - - - Response 2 - . - Response 3 . . . Response 4

DPO | Reference comparison | $\alpha=0.99$, $\lambda=0.8$, $\beta=0.8$
Same prompt in each row; identical response colors and 0-1 axes




DPO | Feedback comparison | $\alpha=0.99$, $\lambda=0.8$, $\beta=0.8$
Same prompt in each row; identical response colors and 0-1 axes




DPO | Relative-sequence entropy | $\alpha=0.99$, $\lambda=0.8$, $\beta=0.8$
 $H(\text{softmax}(s(t)-s(0)))$, natural logarithm; not raw panel entropy


Prompt 54




Prompt 251




Prompt 612




Prompt 737



Prompt 867









Prompt 945









— Ordinary — Reference: $\nu=0.99$ — Feedback: $\kappa=1$
— Reference: $\nu=0.495$ — Feedback: $\kappa=0.5$

DPO | Ordinary | $\alpha=0.9$, $\lambda=0.5$, $\beta=0.8$
Measured panel policy π : four fixed responses, all six prompts









— Response 1 - - - Response 2 - . - Response 3 . . . Response 4

DPO | Reference: nu=0.45 | alpha=0.9, lambda=0.5, beta=0.8
Measured panel policy pi: four fixed responses, all six prompts









— Response 1 - - - Response 2 - . - Response 3 . . . Response 4

DPO | Reference: nu=0.9 | alpha=0.9, lambda=0.5, beta=0.8
Measured panel policy pi: four fixed responses, all six prompts









— Response 1 - - Response 2 - . Response 3 . . . Response 4

DPO | Feedback: kappa=0.5 | alpha=0.9, lambda=0.5, beta=0.8
Measured panel policy pi: four fixed responses, all six prompts




— Response 1 - - - Response 2 - . - Response 3 . . . Response 4

DPO | Feedback: kappa=1 | alpha=0.9, lambda=0.5, beta=0.8
Measured panel policy pi: four fixed responses, all six prompts



— Response 1 - - - Response 2 - . - Response 3 . . . Response 4

DPO | Reference comparison | $\alpha=0.9$, $\lambda=0.5$, $\beta=0.8$
Same prompt in each row; identical response colors and 0-1 axes





















DPO | Feedback comparison | $\alpha=0.9$, $\lambda=0.5$, $\beta=0.8$
Same prompt in each row; identical response colors and 0-1 axes

Ordinary







Feedback: $\kappa=0.5$

Feedback: $\kappa=1$









Response 1 Response 2 Response 3 Response 4

DPO | Relative-sequence entropy | alpha=0.9, lambda=0.5, beta=0.8
 $H(\text{softmax}(s(t)-s(0)))$, natural logarithm; not raw panel entropy









— Ordinary — Reference: nu=0.9 — Feedback: kappa=1
— Reference: nu=0.45 — Feedback: kappa=0.5

DPO | Ordinary | $\alpha=0.9$, $\lambda=0.8$, $\beta=1.2$
Measured panel policy π : four fixed responses, all six prompts









— Response 1 - - - Response 2 - . - Response 3 . . . Response 4

DPO | Reference: nu=0.45 | alpha=0.9, lambda=0.8, beta=1.2
Measured panel policy pi: four fixed responses, all six prompts









— Response 1 - - - Response 2 - . - Response 3 . . . Response 4

DPO | Reference: nu=0.9 | alpha=0.9, lambda=0.8, beta=1.2
Measured panel policy pi: four fixed responses, all six prompts









— Response 1 - - - Response 2 - . - Response 3 . . . Response 4

DPO | Feedback: kappa=0.5 | alpha=0.9, lambda=0.8, beta=1.2
Measured panel policy pi: four fixed responses, all six prompts




— Response 1 - - - Response 2 - . - Response 3 . . . Response 4

DPO | Feedback: kappa=1 | alpha=0.9, lambda=0.8, beta=1.2
Measured panel policy pi: four fixed responses, all six prompts




— Response 1 - - - Response 2 - . - Response 3 . . . Response 4

DPO | Reference comparison | $\alpha=0.9$, $\lambda=0.8$, $\beta=1.2$
Same prompt in each row; identical response colors and 0-1 axes



DPO | Feedback comparison | $\alpha=0.9$, $\lambda=0.8$, $\beta=1.2$
Same prompt in each row; identical response colors and 0-1 axes



DPO | Relative-sequence entropy | alpha=0.9, lambda=0.8, beta=1.2
 $H(\text{softmax}(s(t)-s(0)))$, natural logarithm; not raw panel entropy

