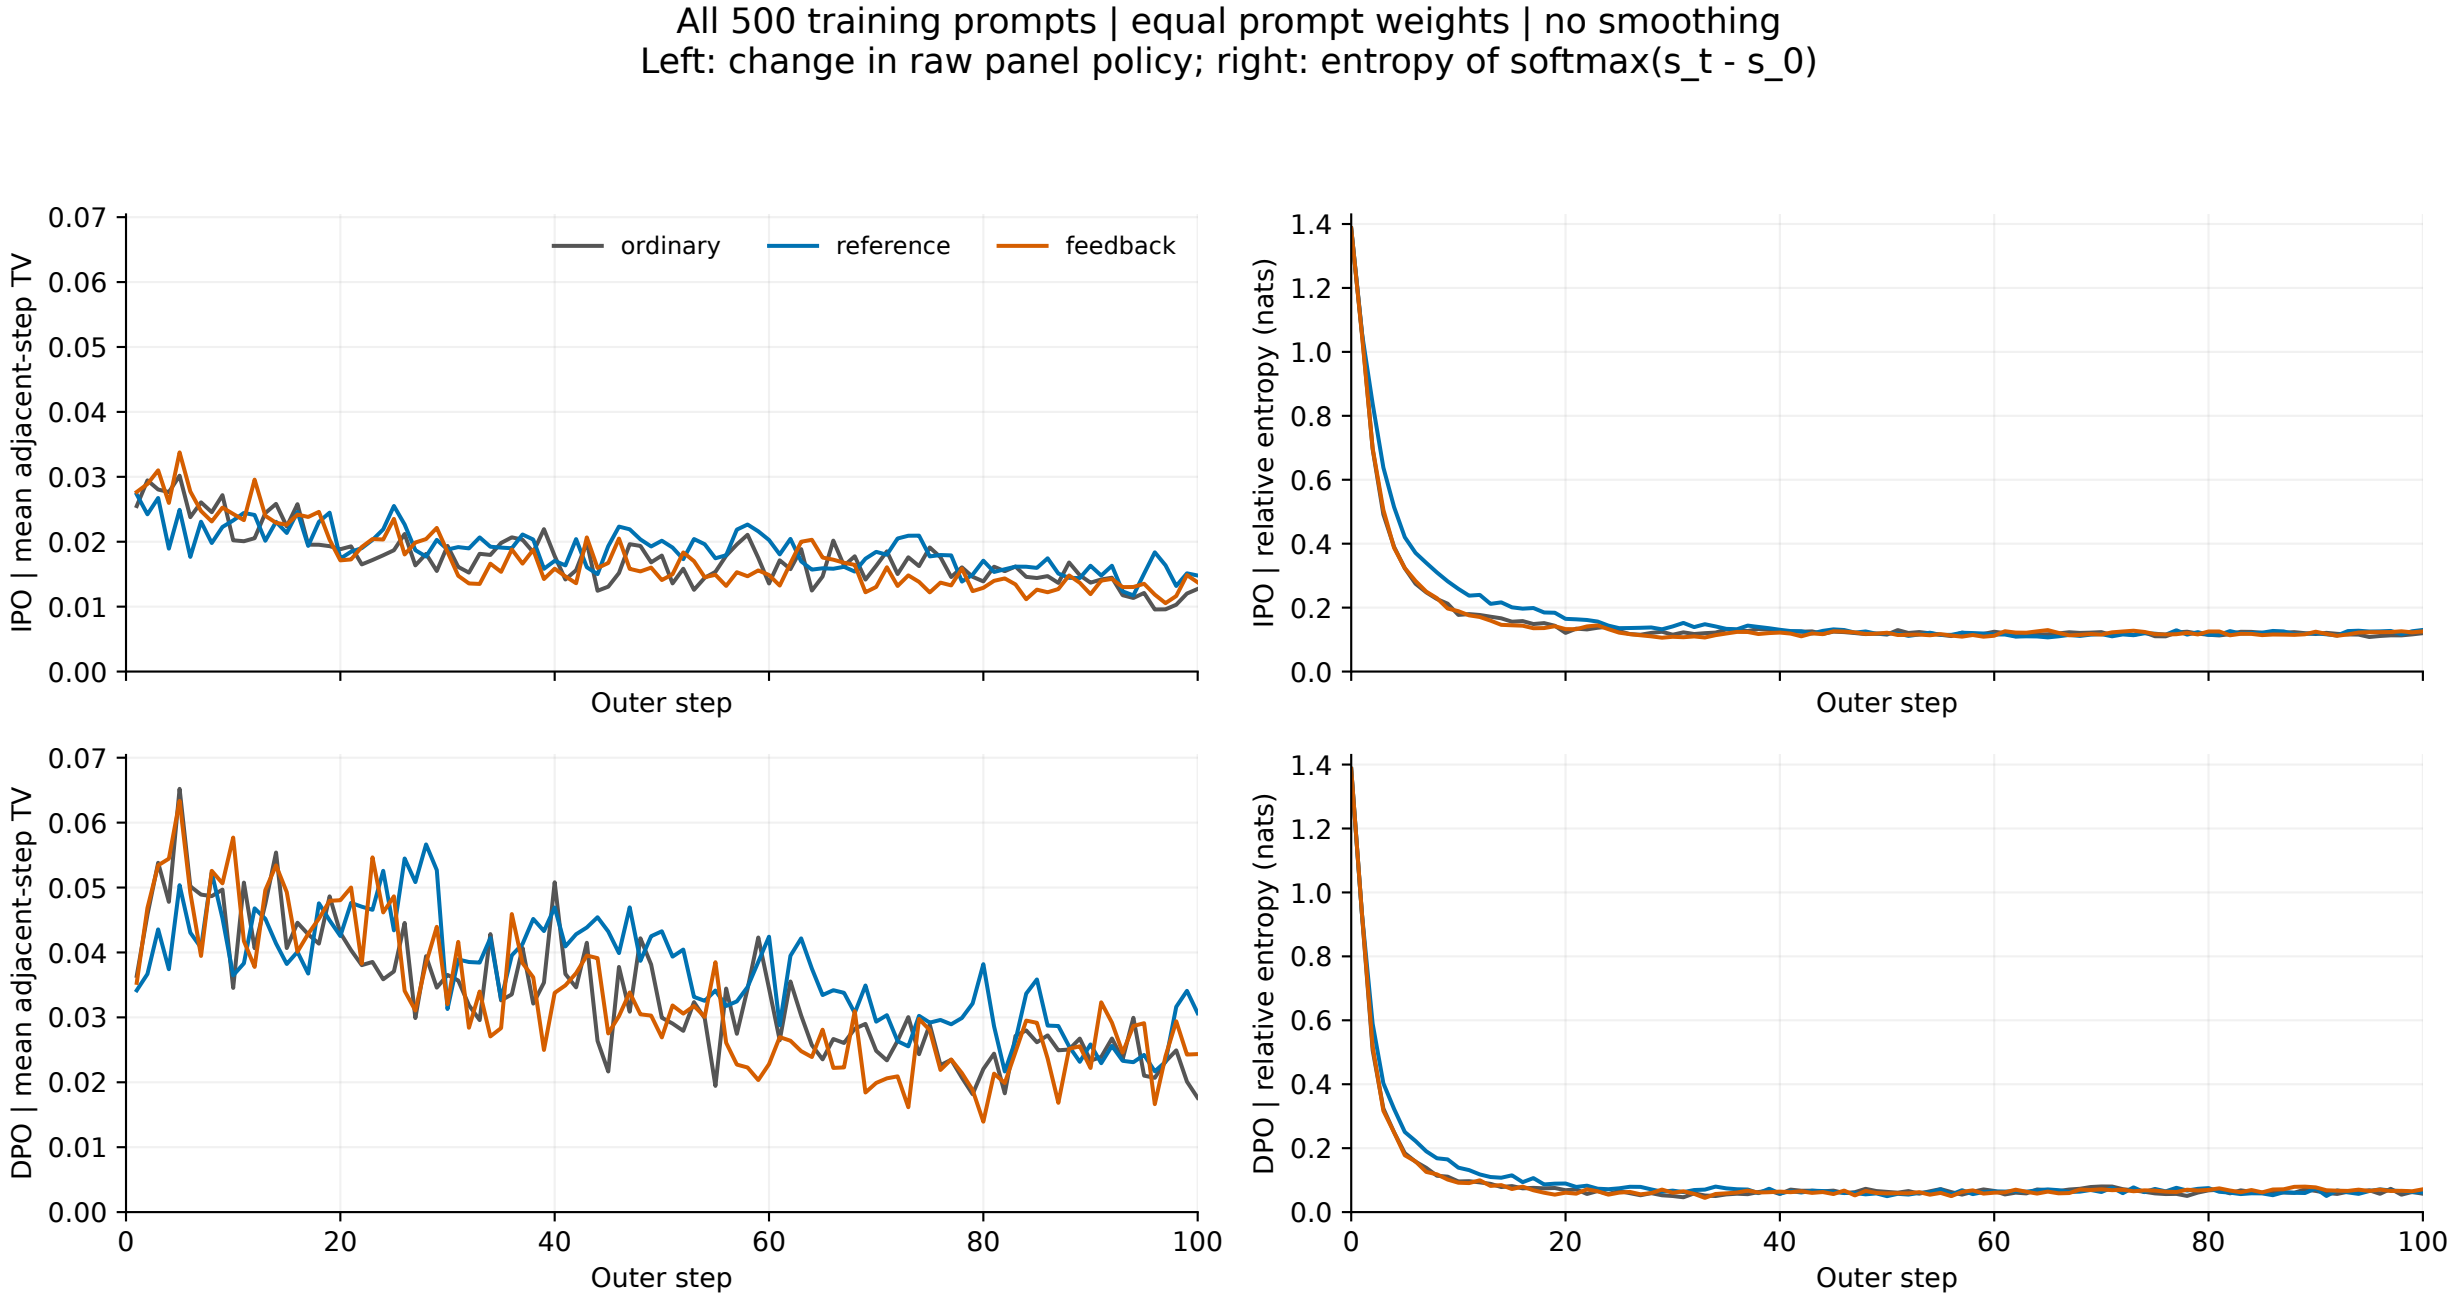
All 500 training prompts | equal prompt weights | no smoothing
Left: change in raw panel policy; right: entropy of softmax(s_t - s_0)
ordinary
reference
feedback
IPO | mean adjacent-step TV
IPO | relative entropy (nats)
DPO | mean adjacent-step TV
DPO | relative entropy (nats)
Outer step
0.00
0.01
0.02
0.03
0.04
0.05
0.06
0.07
0.0
0.2
0.4
0.6
0.8
1.0
1.2
1.4
0
20
40
60
80
100